A Lie Group Approach to Riemannian Batch Normalization

Ziheng Chen, Yue Song, Yunmei Liu, and Nicu Sebe

Code Paper Homepage

UNIVERSITÀ DI TRENTO · UNIVERSITY OF LOUISVILLE · ICLR International Conference On Learning Representations

Motivation

Euclidean Batch Normalization:

Facilitating network training by controlling mean and variance

$$\forall i \leq N, x_i \leftarrow \gamma \frac{x_i - \mu_b}{\sqrt{v_b^2 + \epsilon}} + \beta$$

Existing Riemannian Batch Normalization:

Fails to normalize statistics in a general manner

Table 2: Summary of some representative RBN methods.

Methods	Involved Statistics	Controllable Mean	Controllable Variance	Application Scenarios
SPDBN (Brooks et al., 2019b)	Mean	✓	N/A	SPD manifolds under AIM
SPDBN (Kobler et al., 2022b)	Mean+Variance	✓	✓	SPD manifolds under AIM
Chakraborty (2020, Algs. 1-2)	Mean+Variance	✗	✗	Riemannian homogeneous space
Chakraborty (2020, Algs. 3-4)	Mean+Variance	✓	✓	A certain Lie group structure and distance
RBN (Lou et al., 2020, Alg. 2)	Mean+Variance	✗	✗	Geodesically complete manifolds
Ours	Mean+Variance	✓	✓	General Lie groups

Prelimianries

Lie Groups and Pullback

Definition 2.1 (Lie Groups). A manifold is a Lie group, if it forms a group with a group operation $\odot$ such that $m(x, y) \mapsto x \odot y$ and $i(x) \mapsto x_{\odot}^{-1}$ are both smooth, where $x_{\odot}^{-1}$ is the group inverse.

Definition 2.2 (Left-invariance). A Riemannian metric g over a Lie group $\{G, \odot\}$ is left-invariant, if for any $x, y \in G$ and $V_1, V_2 \in T_x\mathcal{M}$,

$$g_y(V_1, V_2) = g_{L_x(y)}(L_{x*,y}(V_1), L_{x*,y}(V_2)), \tag{1}$$

Definition 2.3 (Pullback Metrics). Suppose $\mathcal{M}_1, \mathcal{M}_2$ are smooth manifolds, g is a Riemannian metric on $\mathcal{M}_2$, and $f : \mathcal{M}_1 \to \mathcal{M}_2$ is smooth. Then the pullback of g by f is defined point-wisely,

$$(f^*g)_p(V, W) = g_{f(p)}(f_{*,p}(V), f_{*,p}(W)), \tag{2}$$

where $p \in \mathcal{M}$, $f_{*,p}(\cdot)$ is the differential map of f at p, and $V, W \in T_p\mathcal{M}$. If f^*g is positive definite, it is a Riemannian metric on $\mathcal{M}_1$, called the pullback metric defined by f.

Geometries on the SPD Manifold

Table 1: Lie group structures and the associated Riemannian operators on SPD manifolds.

Metric	(α, β)-LEM	(α, β)-AIM	LCM
$g_P(V, W)$	$\langle \mathrm{mlog}_{*,P}(V), \mathrm{mlog}_{*,P}(W)\rangle^{(\alpha,\beta)}$	$\langle P^{-1}V, WP^{-1}\rangle^{(\alpha,\beta)}$	$\sum_{i>j} V_{ij}W_{ij} + \sum_{j=1}^{n} V_{jj}W_{jj}L_{jj}^{-2}$
$\mathrm{d}(P, Q)$	$\lVert \mathrm{mlog}(P) - \mathrm{mlog}(Q)\rVert^{(\alpha,\beta)}$	$\left\lVert \mathrm{mlog}\left(Q^{-\frac{1}{2}}PQ^{-\frac{1}{2}}\right)\right\rVert^{(\alpha,\beta)}$	$\lVert \psi_{\mathrm{LC}} \circ \mathrm{Chol}(P) - \psi_{\mathrm{LC}} \circ \mathrm{Chol}(Q)\rVert_{\mathrm{F}}$
$Q \odot P$	$\mathrm{mexp}(\mathrm{mlog}(P) + \mathrm{mlog}(Q))$	$KPK^\top$	$\mathrm{Chol}^{-1}(\lfloor L + K \rfloor + \mathbb{K}\mathbb{L})$
$\mathrm{FM}(\{P_i\})$	$\mathrm{mexp}\left(\frac{1}{n}\sum_i \mathrm{mlog}\, P_i\right)$	Karcher Flow	$\psi_{\mathrm{LC}}^{-1}\left(\frac{1}{n}\sum_i \psi_{\mathrm{LC}}(P_i)\right)$
$\mathrm{Log}_P Q$	$(\mathrm{mlog}_{*,P})^{-1}[\mathrm{mlog}(Q) - \mathrm{mlog}(P)]$	$P^{\frac{1}{2}}\,\mathrm{mlog}\left(P^{-\frac{1}{2}}QP^{-\frac{1}{2}}\right)P^{\frac{1}{2}}$	$(\mathrm{Chol}^{-1})_{*,L}[\lfloor K \rfloor - \lfloor L \rfloor + \mathbb{L}\,\mathrm{Dlog}(\mathbb{L}^{-1}\mathbb{K})]$
Invariance	Bi-invariance	Left-invariance	Bi-invariance

Geometry on the Rotation Matrix

Table 8: The associated Riemannian operators on Rotation matrices.

Operators	$\mathrm{d}^2(R, S)$	$\mathrm{Log}_I R$	$\mathrm{Exp}_I(A)$	$\gamma_{(R,S)}(t)$	FM
Expression	$\left\lVert \mathrm{mlog}\left(R^\top S\right)\right\rVert_{\mathrm{F}}^2$	$\mathrm{mlog}(R)$	$\mathrm{mexp}(A)$	$R\,\mathrm{mexp}(t\,\mathrm{mlog}(R^\top S))$	Manton (2004, Alg. 1)

LieBN on General Lie Groups

Riemannian Gaussian: $p(X \mid M, \sigma^2) = k(\sigma)\exp\left(-\frac{\mathrm{d}(X, M)^2}{2\sigma^2}\right)$

Key operations:

Centering to the neutral element E: $\forall i \leq N, \bar{P}_i \leftarrow L_{M_{\odot}^{-1}}(P_i)$,

Scaling the dispersion: $\forall i \leq N, \hat{P}_i \leftarrow \mathrm{Exp}_E\left[\frac{s}{\sqrt{v^2+\epsilon}}\mathrm{Log}_E(\bar{P}_i)\right]$,

Biasing towards parameter $B \in \mathcal{M}$: $\forall i \leq N, \tilde{P}_i \leftarrow L_B(\hat{P}_i)$,

Proposition 4.1 (Population). [↓] *Given a random point X over $\mathcal{M}$, and the Gaussian distribution $\mathcal{N}(M, v^2)$ defined in Eq. (12), we have the following properties for the population statistics:*

1. *(MLE of M) Given $\{P_{i\ldots N} \in \mathcal{M}\}$ i.i.d. sampled from $\mathcal{N}(M, v^2)$, the maximum likelihood estimator (MLE) of M is the sample Fréchet mean.*
2. *(Homogeneity) Given $X \sim \mathcal{N}(M, v^2)$ and $B \in \mathcal{M}$, $L_B(X) \sim \mathcal{N}(L_B(M), v^2)$*

Proposition 4.2 (Sample). [↓] *Given N samples $\{P_{i\ldots N} \in \mathcal{M}\}$, denoting $\phi_s(P_i) = \mathrm{Exp}_E[s\,\mathrm{Log}_E(P_i)]$, we have the following properties for the sample statistics:*

Homogeneity of the sample mean: $\mathrm{FM}\{L_B(P_i)\} = L_B(\mathrm{FM}\{P_i\}), \forall B \in \mathcal{M}$, (16)

Controllable dispersion from E: $\sum_{i=1}^{N} w_i\,\mathrm{d}^2(\phi_s(P_i), E) = s^2 \sum_{i=1}^{N} w_i\,\mathrm{d}^2(P_i, E)$, (17)

where $\{w_{1\ldots N}\}$ are weights satisfying a convexity constraint, i.e.,*$\forall i, w_i > 0$ and $\sum_i w_i = 1$.*

Algorithm 1: Lie Group Batch Normalization (LieBN) Algorithm

Input : A batch of activations $\{P_{1\ldots N} \in \mathcal{M}\}$, a small positive constant ϵ, and momentum $\gamma \in [0, 1]$
running mean $M_r = E$, running variance $v_r^2 = 1$,
biasing parameter $B \in \mathcal{M}$, scaling parameter $s \in \mathbb{R}/\{0\}$,

Output : Normalized activations $\{\tilde{P}_{1\ldots N}\}$

if *training* **then**
  Compute batch mean M_b and variance v_b^2 of $\{P_{1\ldots N}\}$;
  Update running statistics $M_r \leftarrow \mathrm{WFM}(\{1-\gamma, \gamma\}, \{M_r, M_b\})$, $v_r^2 \leftarrow (1-\gamma)v_r^2 + \gamma v_b^2$;
end
if *training* **then** $M \leftarrow M_b, v^2 \leftarrow v_b^2$;
else $M \leftarrow M_r, v^2 \leftarrow v_r^2$;
for $i \leftarrow 1$ **to** N **do**
  Centering to the neutral element E: $\bar{P}_i \leftarrow L_{M_{\odot}^{-1}}(P_i)$
  Scaling the dispersion: $\hat{P}_i \leftarrow \mathrm{Exp}_E\left[\frac{s}{\sqrt{v^2+\epsilon}}\mathrm{Log}_E(\bar{P}_i)\right]$
  Biasing towards parameter B: $\tilde{P}_i \leftarrow L_B(\hat{P}_i)$
end

Natural Generalization of the EBN

Proposition D.1. *The LieBN algorithm presented in Alg. 1 is equivalent to the standard Euclidean BN when $\mathcal{M} = \mathbb{R}^n$, both during the training and testing phases.*

Gaussian Preservation

$$\phi_s(P) = \mathrm{Exp}_E[s\,\mathrm{Log}_E(P)]$$

Lemma C.1. *Given a random point X distributed over $\mathcal{M}$ with P.D.F. p_X, the P.D.F of $Y = \phi_s(X)$ is given by:*

$$p_Y(Q) = \Delta p_X(\phi_s^{(-1)}(Q)). \tag{26}$$

where $\Delta = \frac{|\phi_{s*}^{-1}|}{\left|L_{\phi_s^{-1}(Q)\odot Q^{-1}*}\right|}$. Here $|\cdot|$ denotes the determinant, and ϕ_{s*}^{-1} and $L_{\phi_s^{-1}(Q)\odot Q^{-1}*}$ are the differentials.

Corollary C.2. *Following the notations in Lem. C.1, if $\Delta = c$ is a constant and $X \sim \mathcal{N}(E, \sigma^2)$, then Y also follows a Gaussian distribution,* i.e.,*$Y \sim \mathcal{N}(E, s^2\sigma^2)$*

Corollary C.4. *Given a Lie group $\mathcal{M}$ pulled back from the Euclidean space, and a random point $X \sim \mathcal{N}(E, \sigma^2)$ over $\mathcal{M}$, $Y = \phi_s(X) \sim \mathcal{N}(E, s^2\sigma^2)$*

LEM, LCM and their variants:

$$\mathcal{N}(M, \sigma^2) \to \mathcal{N}(E, \sigma^2) \to \mathcal{N}(E, s^2) \to \mathcal{N}(B, s^2)$$

LieBN on the Deformed SPD Lie Groups

Power deformation: $\tilde{g} = \frac{1}{\theta^2}\mathrm{Pow}_\theta^* g$,

Table 3: Key operators in calculating LieBN on SPD manifolds.

	Metric	(θ, α, β)-AIM	(α, β)-LEM	θ-LCM
	Pullback Map	P_θ	mlog	$\mathrm{P}_\theta \circ \psi_{\mathrm{LC}}$
	Codomain	$\{\mathcal{S}_{++}^n, \odot^{\mathrm{AI}}, \frac{1}{\theta^2}g^{(\alpha,\beta)\text{-AI}}\}$	$\{\mathcal{S}^n, \langle\cdot,\cdot\rangle^{(\alpha,\beta)}\}$	$\{\mathcal{L}^n, \frac{1}{\theta^2}\langle\cdot,\cdot\rangle\}$
Riemannian and Lie group operators in the codomain	$L_Q(P)$	$KPK^\top$	$P+Q$	$P+Q$
	$L_{Q_{\odot}^{-1}}(P)$	$K^{-1}PK^{-\top}$	$P-Q$	$P-Q$
	$\mathrm{Exp}_E[s\,\mathrm{Log}_E(P)]$	P^s	sP	sP
	FM	Karcher Flow	Arithmetic average	Arithmetic average
	$\mathrm{WFM}(\{\gamma, 1-\gamma\}, \{P_1, P_2\})$	$P_2^{\frac{1}{2}}\left(P_2^{-\frac{1}{2}}P_1P_2^{-\frac{1}{2}}\right)^\gamma P_2^{\frac{1}{2}}$	Arithmetic weighted average	Arithmetic weighted average

Proposition 5.1 (Deformation). [↓] *(θ, α, β)-LEM is equal to (α, β)-LEM. θ-LCM interpolates between $\tilde{g}$-LEM ($\theta = 0$) and LCM ($\theta = 1$), with $\tilde{g}$-LEM defined as*

$$\langle V, W\rangle_P = \tilde{g}(\mathrm{mlog}_{*,P}(V), \mathrm{mlog}_{*,P}(W)), \forall P \in \mathcal{S}_{++}^n, \forall V, W \in T_P\mathcal{S}_{++}^n, \tag{18}$$

where $\tilde{g}(V_1, V_2) = \frac{1}{2}\langle V_1, V_2\rangle - \frac{1}{4}\langle \mathbb{D}(V_1), \mathbb{D}(V_2)\rangle$, $\mathbb{D}(V_i)$ is a diagonal matrix consisting of the diagonal elements of V_i, and $\mathrm{mlog}_{*,P}$ is the differential map at P.

Proposition 5.2 (Invariance). [↓] *(θ, α, β)-AIM is left-invariant w.r.t. $\odot^{\theta\text{-AI}}$, while θ-LCM is bi-invariant w.r.t. $\odot^{\theta\text{-LC}}$.*

Experiments on the SPD and Rotation Matrices

Table 4: 10-fold average results of SPDNet with and without SPDBN or LieBN on the Radar, HDM05, and FPHA datasets. For simplicity, LieBN-Metric-(θ) is abbreviated as Metric-(θ). For the LieBN under each metric, if the LieBN induced by the standard metric ($\theta = 1$) is not saturated, we report the LieBN under the deformed metric in the rightmost columns of the table.

(a) Radar dataset.

Method	SPDNet	SPDNetBN	AIM-(1)	LEM-(1)	LCM-(1)	LCM-(-0.5)
Fit Time (s)	0.98	1.56	1.62	1.28	1.11	1.43
Mean±STD	93.25±1.10	94.85±0.99	**95.47±0.90**	94.89±1.04	93.52±1.07	94.80±0.71
Max	94.4	96.13	96.27	**96.8**	95.2	95.73

(b) HDM05 and FPHA datasets.

Method		SPDNet	SPDNetBN	AIM-(1)	LEM-(1)	LCM-(1)	AIM-(1.5)	LCM-(0.5)
HDM05	Fit Time (s)	0.57	0.97	1.14	0.87	0.66	1.46	1.01
	Mean±STD	59.13±0.67	66.72±0.52	67.79±0.65	65.05±0.63	66.68±0.71	68.16±0.68	**70.84±0.92**
	Max	60.34	67.66	68.75	66.05	68.52	69.25	**72.27**
FPHA	Fit Time (s)	0.32	0.62	0.80	0.55	0.39	1.03	0.65
	Mean±STD	85.59±0.72	89.33±0.49	89.70±0.51	86.56±0.79	77.64±1.00	**90.39±0.66**	86.33±0.43
	Max	86	90.17	90.5	87.83	79	**92.17**	87

Table 5: Cross-validation results of TSMNet with SPDDSMBN and DSMLieBN on the Hinss dataset. For simplicity, DSMLieBN-Metric-(θ) is abbreviated as Metric-(θ). For the DSMLieBN under each metric, if the DSMLieBN induced by the standard metric ($\theta = 1$) is not saturated, we report the DSMLieBN under the deformed metric at the bottom rows of the table.

(a) Inter-session classification

Method	Fit Time (s)	Mean±STD
SPDDSMBN	0.16	54.12±9.87
AIM-(1)	0.16	**55.10±7.61**
LEM-(1)	0.13	54.95±10.09
LCM-(1)	0.10	51.54±6.88
LCM-(0.5)	0.15	53.11±5.65

(b) Inter-subject classification

Method	Fit Time (s)	Mean±STD
SPDDSMBN	7.74	50.10±8.08
AIM-(1)	6.94	50.04±8.01
LEM-(1)	4.71	50.95±6.40
LCM-(1)	3.59	51.86±4.53
AIM-(-0.5)	8.71	**53.97±8.78**

Table 9: Results of LieNet with or without LieBN on the G3D dataset.

Methods	G3D	
	Mean±STD	Max
LieNet	87.91±0.90	89.73
LieNetLieBN	88.88±1.62	90.67